




Runtime feedback closes the loop: the schedule solved on measured board costs executes on time where the isolated-profile one slides



Same spec (wh_chain120_solve_h200), same board, same solver.

The isolated-profile schedule assumes YOLO is cheap; on the board it runs long, downstream harts wait, and the timeline slides ~168 ms behind a 200 ms plan with harts only 47% busy.

The board-calibrated schedule (per-dispatch YOLO cost measured under load) keeps harts 68% busy and holds every window.

A global service-time multiplier fitted from the isolated run does not bridge them: the bad schedule under-contains YOLO, so the fit under-weights it and re-solves the same shape.